

Lab 3
(2h40 hours, evaluated)

Neural Networks
Implementing a Neural Network with Keras

Introduction

In this lab, we will see how to build a neural network with the Keras framework in Python. As usual, we give an incomplete python program. Some instructions will be used as they are and you do not need to change them. We will comment some of them but we do not need to understand them all. We will use the dataset MNIST, one of the most used datasets in machine learning research. This dataset consists of 70000 images of handwritten digits (of size 28x28).

So first, you need to download the files lab3_skeleton.py and lab3_3_skeleton.py available at:

[/net/ens/DeepLearning/DLCV2020/lab3_skeleton.py](#)

[/net/ens/DeepLearning/DLCV2020/lab3_3_skeleton.py](#)

or http://dept-info.labri.fr/~mansenca/DLCV2020/lab3_skeleton.py

http://dept-info.labri.fr/~mansenca/DLCV2020/lab3_3_skeleton.py

You will need to activate the work environment with:

source /net/ens/DeepLearning/python3/tensorflow2/bin/activate

Basics:

Recall that a neural network (NN for short) consists of a set of layers disposed linearly, and each layer is a set of (artificial) neurons. The model is simplified so that signals can only circulate from the bottom layer to the top layer. Each neuron in the k -th layer of the neural network is connected to all the neurons in the $(k - 1)$ -th layer, and the neurons in a given layer are all independent from each other.

A neural network consists of the following components

- An **input layer**, x
- An arbitrary amount of **hidden layers**
- An **output layer**, y
- A set of **weights** and **biases** between each layer, w and b
- A choice of **activation function** for each hidden layer, σ .

1. Single Neuron:

As a first exercise, we will consider a single neuron with $28 \times 28 = 784$ inputs and a single sigmoid unit generating the output. Our neuron will simply learn to distinguish between the digit 0 and the other digits (in the MNIST dataset).

1. Complete the file with the model definition in keras
2. Train your network and measure its accuracy.
3. Try various batch sizes.

2. A Neural Network with One Hidden Layer:

In this section, we will add a hidden layer with 64 units.

1. Modify your file accordingly to implement this NN.
2. Train your network and measure its accuracy.
3. Try various number of neurons on hidden layer
4. Try various activation functions

3. Multiclass Neural Network:

Now, we will adapt our previous network to recognize all the 10 digits.

1. Modify your file accordingly to implement this NN.
2. Train your network and measure its accuracy.
3. Try various optimizers
4. Search the best network architecture that maximize accuracy

You must send an archive with your code.

This archive must be named [LASTNAME_firstname_Lab3_DLCV.tar.bz2](#) and sent to: boris.mansencal@labri.fr with DLCV, your group day, name and lab number in the e-mail subject: for example “[DLCV-thursday: DOE_John_Lab3](#)”.

In your report, please note the loss value for 100 epochs for each method, plot the loss over epochs, compute and discuss the obtained accuracies.

For each method, try to find “good” hyper-parameters (i.e., try to maximize accuracy).